

Class 07: Machine Learning 1


Yufei (A16222438)

Clustering

We will start today's lab with clustering methods, in particular so-called **K-means**. The main function for this in R is `kmeans()`.

Let's try it on some made-up data where we know the answer.

```
# rnorm function draws random points from a normal distribution
# rnorm(number of points, mean, sd). default: z.
x <- rnorm(1000, mean = 3)
hist(x)
```




make 60 points

```
# points from two distributions
tmp <- c(rnorm(30, mean = 3), rnorm(30, mean = -3))
# create a y axis as the reversed list of tmp, build matrix by cbind
tdtmp <- cbind(x=tmp, y=rev(tmp))
#check the data
head (tdtmp)
```

```
      x      y
[1,] 3.407357 -3.3618319
[2,] 2.076548 -3.3447115
[3,] 1.801669 -1.4879422
[4,] 3.927693 -0.9860571
[5,] 0.670952 -2.4202559
[6,] 3.312089 -3.0992685
```

```
#plot the data
plot (tdtmp)
```



K-means clustering of two sets of points of size 30, 30.

```
# identify the two clusters by kmeans(). Iteration = 20 times.
k <- kmeans(tdtmp, centers = 2, nstart = 20)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	2.902282	-3.001785
2	-3.001785	2.902282

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 60.22932 60.22932
(between_SS / total_SS = 89.7 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Q1. How many points are in each cluster?

```
k$size
```

```
[1] 30 30
```

Q2 Cluster membership?

```
k$cluster
```

[illegible]


Q3. Cluster centers?

```
k$centers
```

	x	y
1	2.902282	-3.001785
2	-3.001785	2.902282


Q4. Plot my clustering results?

```
plot(tdtmp, col = k$cluster+2, pch = 16)
```



Q5. Cluster the data again with `kmeans()` into 4 groups and plot the results.

```
kf <- kmeans(tdtmp, centers = 4, nstart = 20)
plot(tdtmp, col = kf$cluster, pch = 16)
```



K-means is very popular mostly because it is fast and relatively straightforward to run and understand. But it has a big limitation in that you need to tell it how many groups (k, or centers) you want. If we don't know k beforehand, we need trial and error to obtain the optimal result with minimal SS and cluster number.

Hierachical clustering

The main function in base R is called `hclust()`. You have to pass it in a “distance matrix” not just your input data.

We can generate a distance matrix with the `dist()` function.

```
hc <- hclust (dist(tdtmp))
hc
```

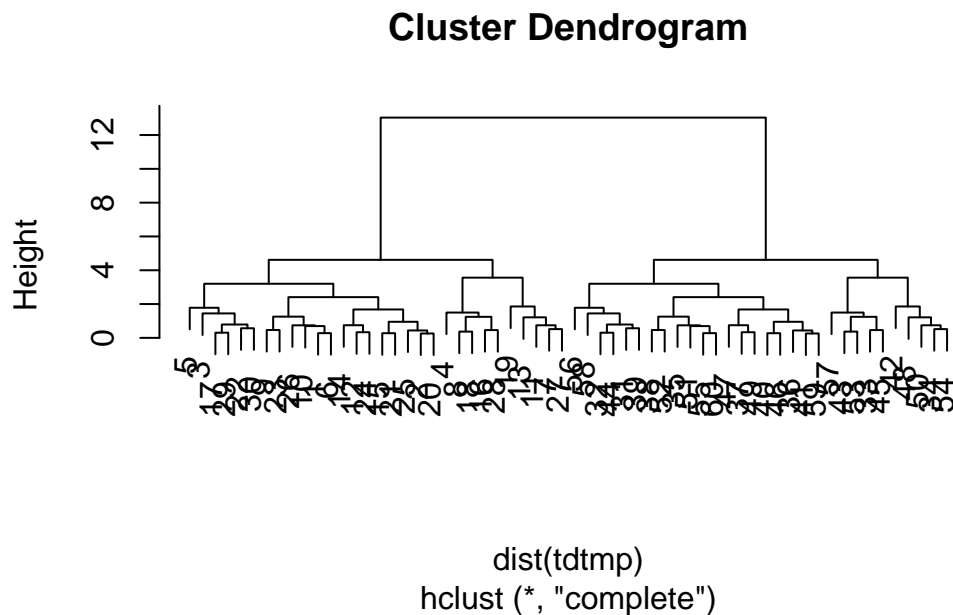
Call:

```
hclust(d = dist(tdtmp))
```

```
Cluster method : complete
Distance       : euclidean
```

Number of objects: 60

```
plot(hc)
```




Interpretation: - It grouped the numbers into two groups: ≤ 30 (left) and > 30 (right). - Height of bars = distance of points. - Determine number of clusters by observing the number of highest bars.

To find clusters (cluster membership vector) from a ‘hclust()’ result, we can “cut” the tree at a certain height that we like

```
# cut the tree on plot  
plot(hc)  
abline(h=8, col = "red")
```

Cluster Dendrogram




```
dist(tdtmp)  
hclust (*, "complete")
```

```
#get membership vector  
grps <- cutree(hc, h=8)  
table(grps)
```

```
grps  
1 2  
30 30
```

Plot the result

```
plot(tdtmp, col=grps+2, pch = 16)
```



PCA: Principal Component Analysis

Read the data consumption in grams (per person, per week) of 17 different types of food-stuff measured and averaged in the four countries of the United Kingdom in 1997.

```
x <- read.csv("https://tinyurl.com/UK-foods")
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 5
```



```
#or
dim(x)
```

```
[1] 17  5
```

```
# Checking data
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

It looks like the row-names here were not set properly as we were expecting 4 columns (one for each of the 4 countries of the UK - not 5 as reported from the `dim()` function)

```
rownames(x) <- x[,1]
food <- x[,-1]
head(x)
```

	X	England	Wales	Scotland	N.Ireland
Cheese	Cheese	105	103	103	66
Carcass_meat	Carcass_meat	245	227	242	267
Other_meat	Other_meat	685	803	750	586
Fish	Fish	147	160	122	93
Fats_and_oils	Fats_and_oils	193	235	184	209
Sugars	Sugars	156	175	147	139

It is dangerous way of processing data as it deletes the first column every time we run this code, until all columns removed.

Instead, we can read the data properly


```
x <- read.csv("https://tinyurl.com/UK-foods", row.names = 1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

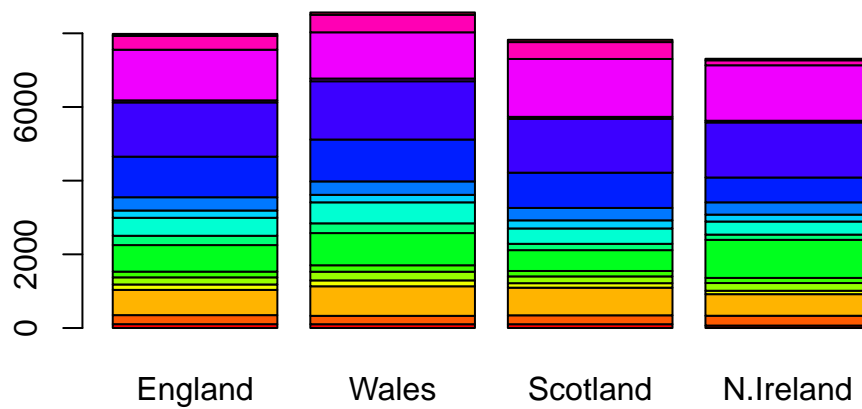
We can do bar plots

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?


```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

A pairs plot can be useful, if we don't have much dimensions

```
pairs(x, col = rainbow(nrow(x)), pch = 16)
```



Interpretation: X and Y axis represent data from corresponding country, color of dot represent food category. Deviation from diagonal represent difference in data. Drawback: time-consuming and confusing to interpret.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

PCA of UK Food Data (17 dimensions)

PCA can help us make sense of these types of datasets. Let's see how it works.

The main function in "base" R is called `prcomp()`. It needs matrix column to be the dimensions of analysis (in this case food type), and row to be group (in this case countries).

```
#transpose x by `t()` function
head(t(x))
```

	Cheese	Carcass_meat	Other_meat	Fish	Fats_and_oils	Sugars
England	105	245	685	147	193	156
Wales	103	227	803	160	235	175
Scotland	103	242	750	122	184	147
N.Ireland	66	267	586	93	209	139

	Fresh_potatoes	Fresh_Veg	Other_Veg	Processed_potatoes	
England	720	253	488		198
Wales	874	265	570		203
Scotland	566	171	418		220
N.Ireland	1033	143	355		187
	Processed_Veg	Fresh_fruit	Cereals	Beverages	Soft_drinks
England	360	1102	1472	57	1374
Wales	365	1137	1582	73	1256
Scotland	337	957	1462	53	1572
N.Ireland	334	674	1494	47	1506
	Alcoholic_drinks	Confectionery			
England	375	54			
Wales	475	64			
Scotland	458	62			
N.Ireland	135	41			

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:


	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Interpretation: PC1 captures 67.44% of variance, and PC2 captures 96.5% variance in data.

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13

```
plot(pca$x[,1],pca$x[,2], col = c("orange", "red", "blue", "darkgreen"), pch = 16)
```



The “loadings” tell us how much the original variables (in our case, food) contribute to the new variable (PC).

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.694538519
Carcass_meat	0.047927628	0.013915823	0.06367111	0.489884628
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.279023718
Fish	-0.084414983	-0.050754947	0.03906481	-0.008483145
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.076097502
Sugars	-0.037620983	-0.043021699	-0.03605745	0.034101334
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	-0.090972715
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	-0.039901917
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.016719075
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	0.030125166
Processed_Veg	-0.036488269	-0.045451802	0.05289191	-0.013969507
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.184072217
Cereals	-0.047702858	-0.212599678	-0.35884921	0.191926714
Beverages	-0.026187756	-0.030560542	-0.04135860	0.004831876
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.103508492
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.316290619

Confectionery -0.029650201 0.005949921 -0.05232164 0.001847469

```
#plot PC1, with graph margin of 2
barplot(pca$rotation[,1], las = 2)
```

